



ARTIFICIAL INTELLIGENCE

2023/2024 Semester 2

Intelligent Agents: Chapter 2

Outline

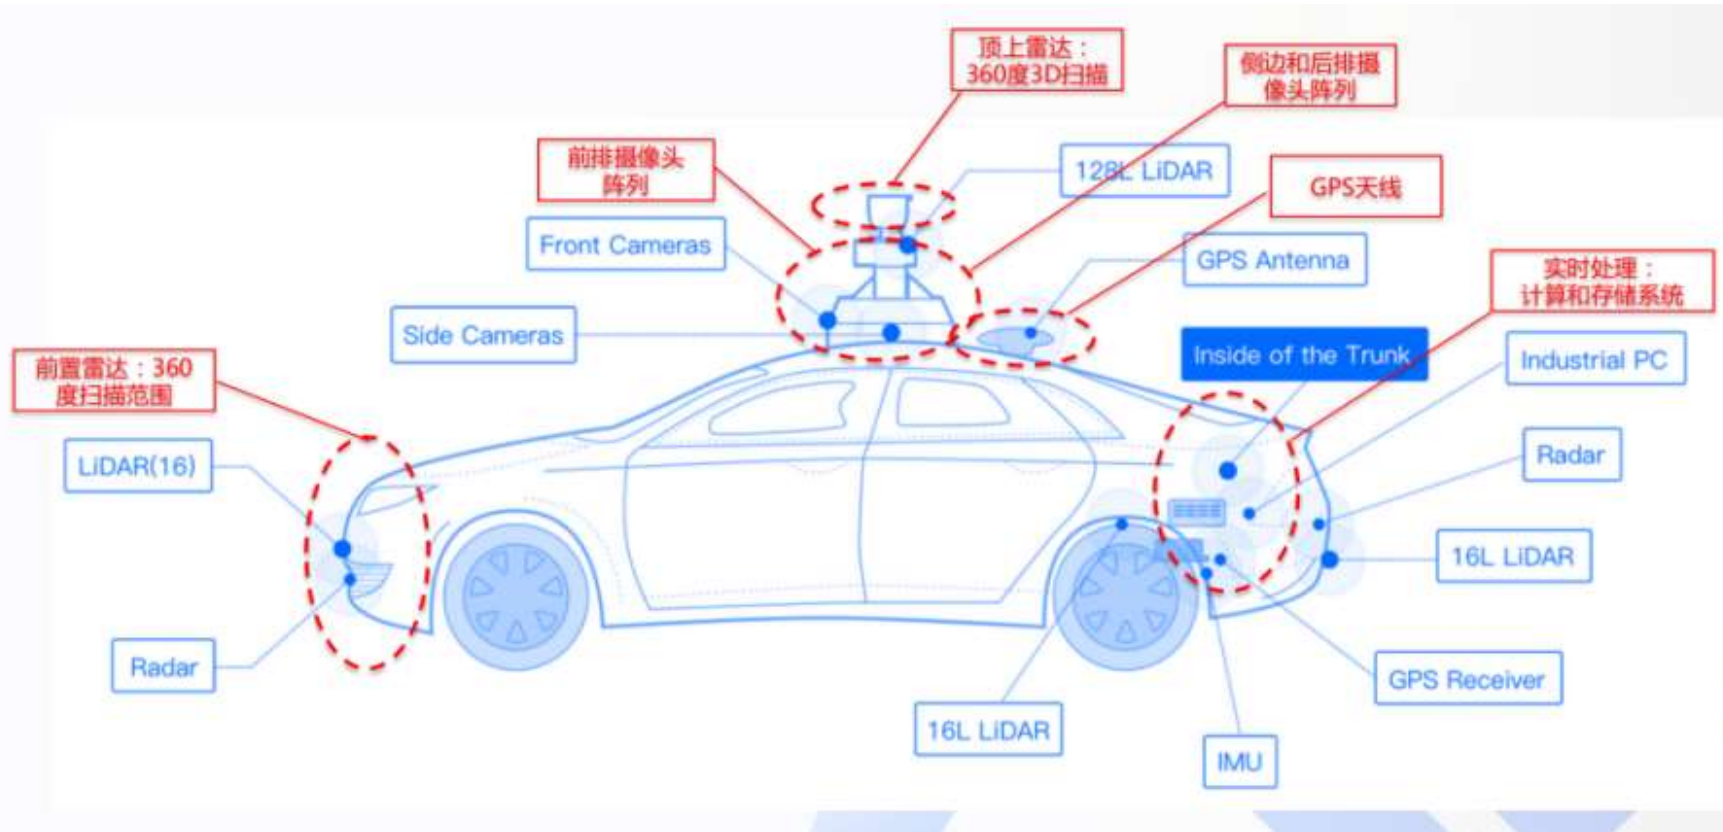
- Agents and environments
- Rationality
- PEAS (**P**erformance measure, **E**nvironment, **A**ctuators, **S**ensors)
- Environment types
- Agent types

Agents


- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**
- **Human agent**: eyes, ears, and other organs for **sensors**; hands, legs, mouth, and other body parts for **actuators**
- **Robotic agent**: cameras and infrared range finders for sensors; various motors for actuators

Agents

- Self-driving Car agent




Agents and environments



- The **agent function** maps from percept histories to actions:
$$[f: \mathcal{P}^* \rightarrow \mathcal{A}]$$
- The **agent program** runs on the physical **architecture** to produce f
- agent = architecture + program

Vacuum-cleaner world



- **Percepts:** location and contents, e.g., [A,Dirty]
- **Actions:** *Left, Right, Suck, NoOp*

Vacuum-cleaner world

Example vacuum agent program:

function Vacuum-Agent([location, status])

returns an **action**

if status = Dirty *then* return **Suck**

else if location = A *then* return **Right**

else if location = B *then* return **Left**

Rational agents

- An agent should strive to "**do the right thing**", based on what it can perceive and the actions it can perform. The right action is the one that will **cause the agent to be most successful**.
- **Performance measure:** An objective criterion for success of an agent's behavior.
- E.g., performance measure
 - amount of dirt cleaned up, amount of time taken
 - amount of electricity consumed, amount of noise generated

Rational agents

- **Rational Agent:** For each possible percept sequence, a rational agent should select an action that is expected
 - to maximize its performance measure,
 - given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.

Rational agents

- **Rationality** is distinct from **omniscience** (all-knowing with infinite knowledge)
- Agents can perform actions in order to modify future percepts so as to obtain useful information (**information gathering**, exploration)
- An agent is **autonomous** if its behavior is determined by its own experience (with ability to **learn and adapt**)


Back to vacuum-cleaner world

- **Percepts:**

Location and status,
e.g., [A, Dirty]

- **Actions:**

Left, Right, Suck, NoOp



function Vacuum-Agent([location, status]) returns an **action**

- *if status = Dirty then return Suck*
- *else if location = A then return Right*
- *else if location = B then return Left*
- Is this agent rational?
 - Depends on performance measure, environment properties

PEAS

- **PEAS**: **P**erformance measure, **E**nvironment, **A**ctuators, **S**ensors
- Must first specify the setting for intelligent agent design
- Consider, e.g., the task of designing an automated taxi driver:
 - Performance measure
 - Environment
 - Actuators
 - Sensors

PEAS

- Must first specify the setting for intelligent agent design
- Consider, e.g., the task of designing an automated taxi driver:
 - **Performance measure:** Safe, fast, legal, comfortable trip, maximize profits
 - **Environment:** Roads, other traffic, pedestrians, customers
 - **Actuators:** Steering wheel, accelerator, brake, signal, horn
 - **Sensors:** Cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboard

PEAS

- Agent: Medical diagnosis system
 - **Performance measure:** Healthy patient, minimize costs, lawsuits
 - **Environment:** Patient, hospital, staff
 - **Actuators:** Screen display (questions, tests, diagnoses, treatments, referrals)
 - **Sensors:** Keyboard (entry of symptoms, patient's answers, examination reports)

PEAS

- Agent: Part-picking robot
 - **Performance measure:** Percentage of parts in correct bins
 - **Environment:** Conveyor belt with parts, bins
 - **Actuators:** Jointed arm and hand
 - **Sensors:** Camera, joint angle sensors

PEAS

- **Agent: Interactive English tutor**
 - **Performance measure:** Maximize student's score on test
 - **Environment:** Set of students
 - **Actuators:** Screen display, Speaker (exercises, suggestions, corrections)
 - **Sensors:** Keyboard

PEAS

- Agent: Spam filter
 - **Performance measure: ?**
 - **Environment :?**
 - **Actuators:?**
 - **Sensors :?**

Environment types

- **Fully observable** (vs. partially observable): An agent's sensors give it access to the complete state of the environment at each point in time.
- **Deterministic** (vs. stochastic): The next state of the environment is completely determined by the current state and the action executed by the agent. (If the environment is deterministic except for the actions of other agents, then the environment is **strategic**)
- **Episodic** (vs. sequential): The agent's experience is divided into atomic "episodes" (each episode consists of the agent perceiving and then performing a single action), and the choice of action in each episode depends only on the episode itself.

Environment types

- **Static** (vs. dynamic): The environment is unchanged while an agent is deliberating. (The environment is **semidynamic** if the environment itself does not change with the passage of time but the agent's performance score does)
- **Discrete** (vs. continuous): A limited number of distinct, clearly defined percepts and actions.
- **Single agent** (vs. multiagent): An agent operating by itself in an environment.

Environment types

| | Chess with a clock | Chess without a clock | Taxi driving |
|------------------|-----------------------|--------------------------|--------------|
| Fully observable | Yes | Yes | No |
| Deterministic | Strategic | Strategic | No |
| Episodic | No | No | No |
| Static | Semi | Yes | No |
| Discrete | Yes | Yes | No |
| Single agent | No | No | No |

- The environment type largely determines the agent design
- The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent

Agent functions and programs

- An agent is completely specified by the agent function mapping percept sequences to actions
- One agent function (or a small equivalence class) should be rational
- **Aim:** find a way to implement the rational agent function concisely

Table-lookup agent

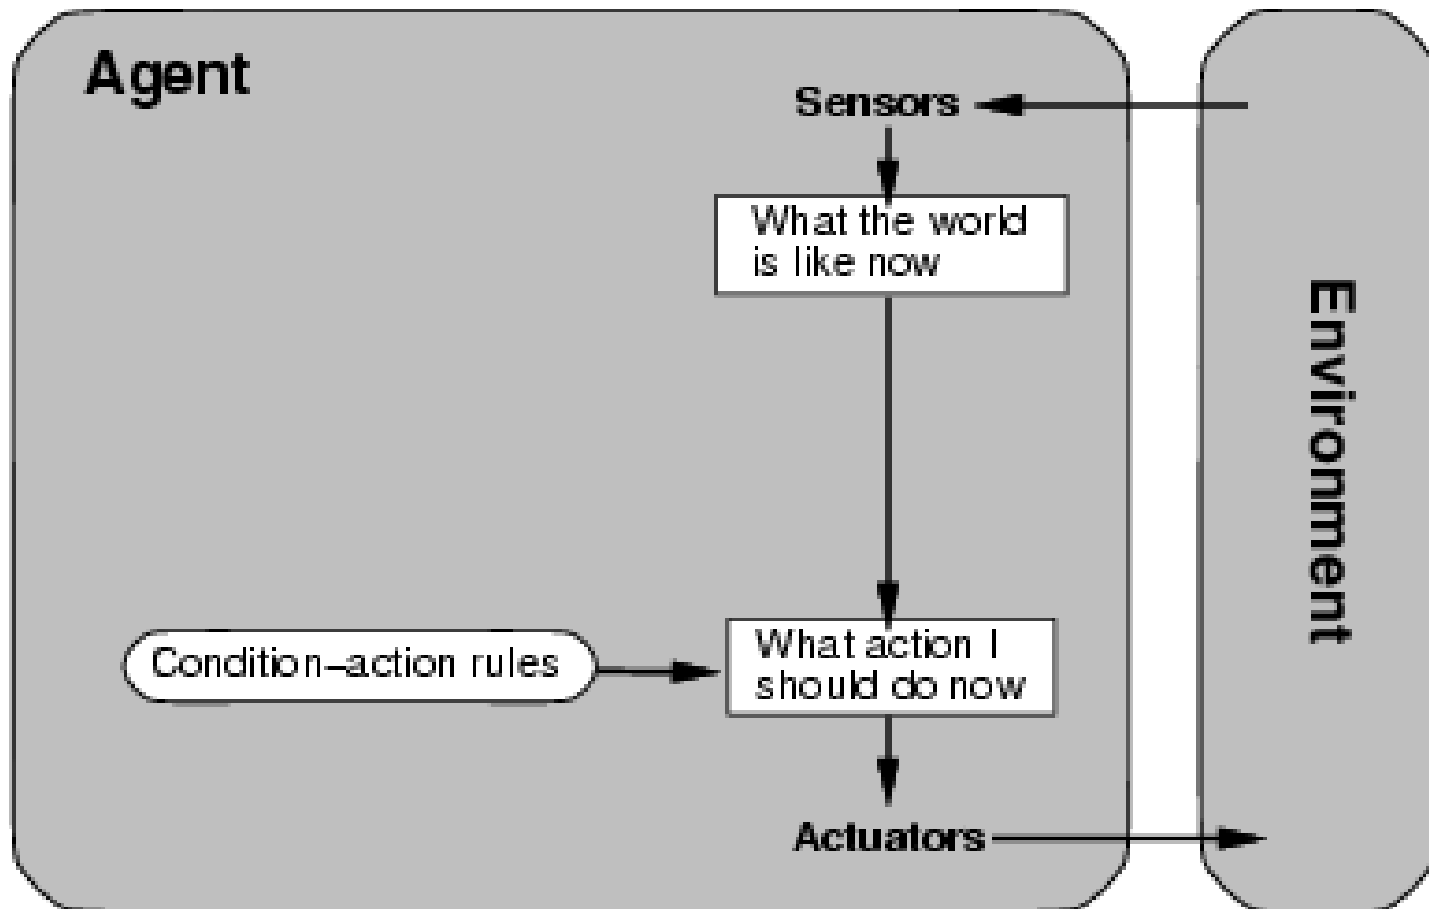
- `\input{algorithms/table-agent-algorithm}`
- Drawbacks:
 - Huge table
 - Take a long time to build the table
 - No autonomy
 - Even with learning, need a long time to learn the table entries

Agent types

- Four basic types in order of increasing generality:
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents

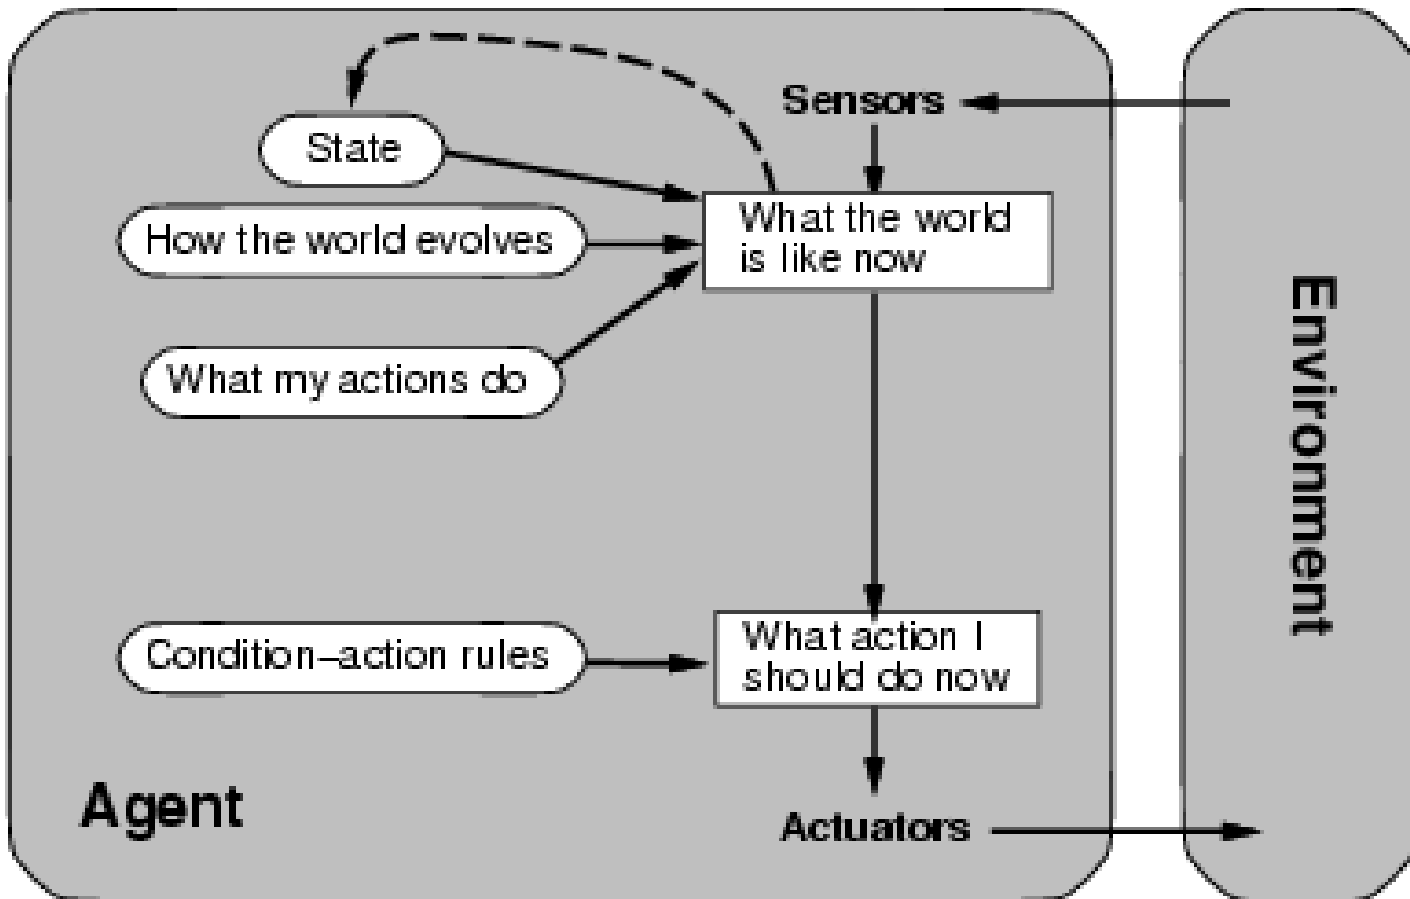
Simple reflex agents

Select action on the basis of current percept, ignoring all past percepts



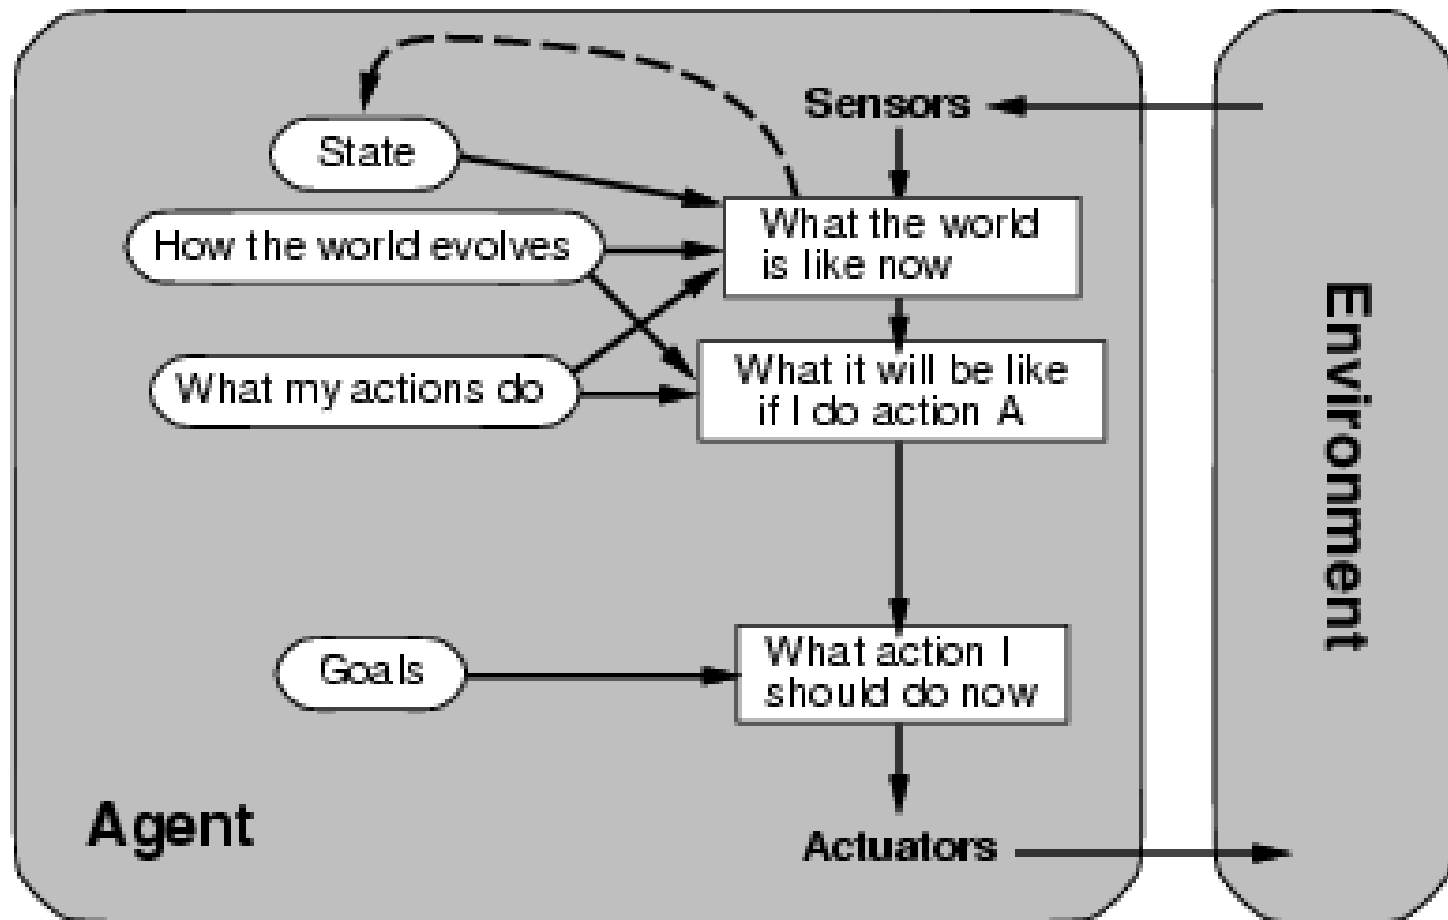
Model-based reflex agents

Maintains internal state that keeps track of aspects of the environment that cannot be currently observed



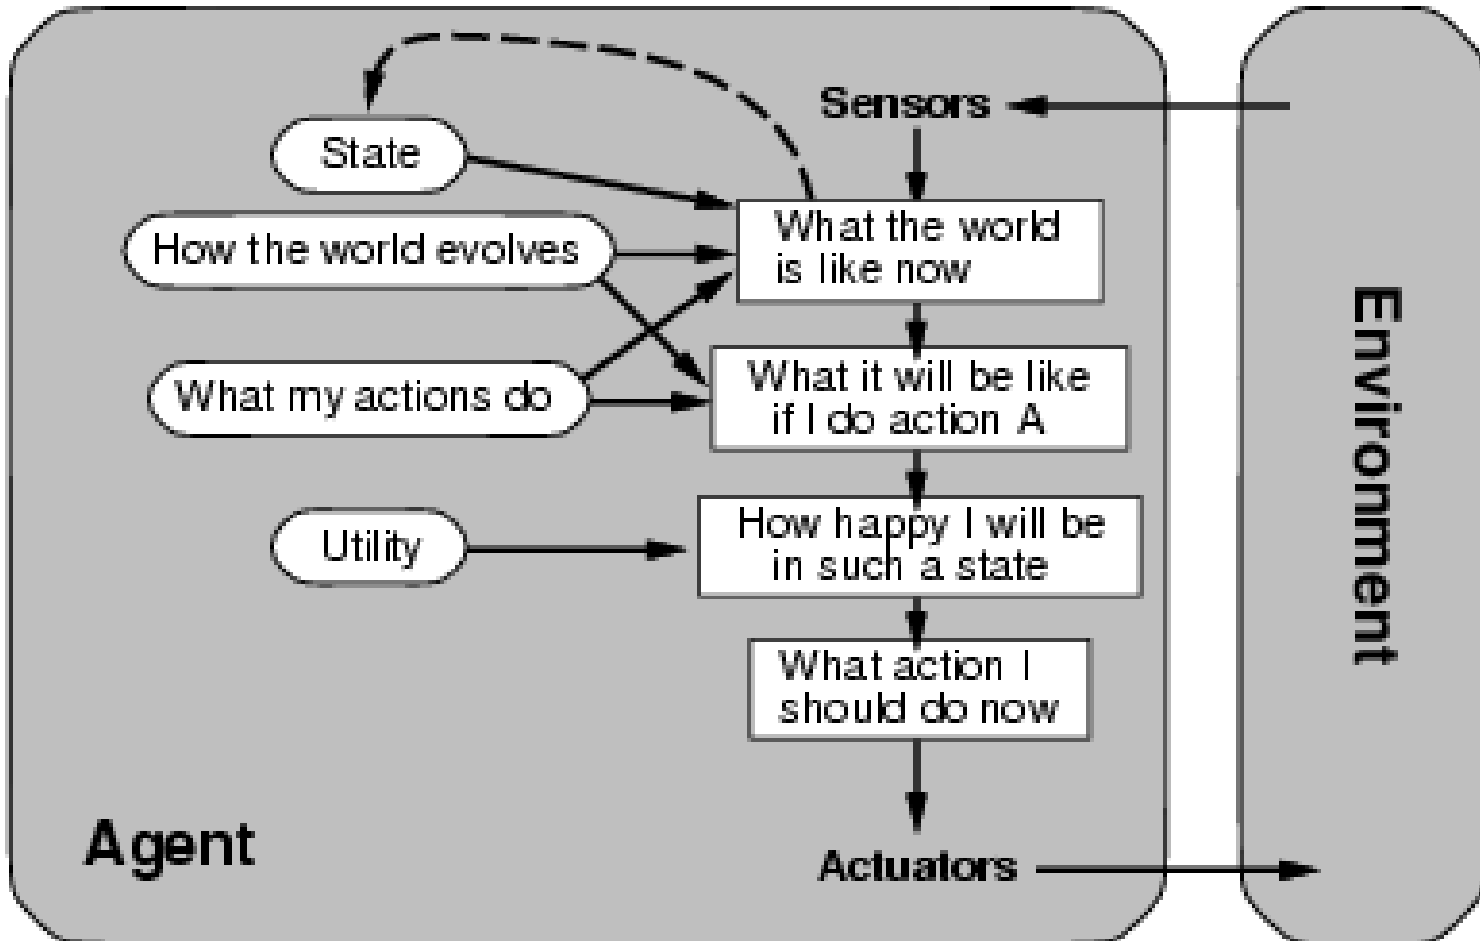
Goal-based agents

The agent uses goal information to select between possible actions in the current state

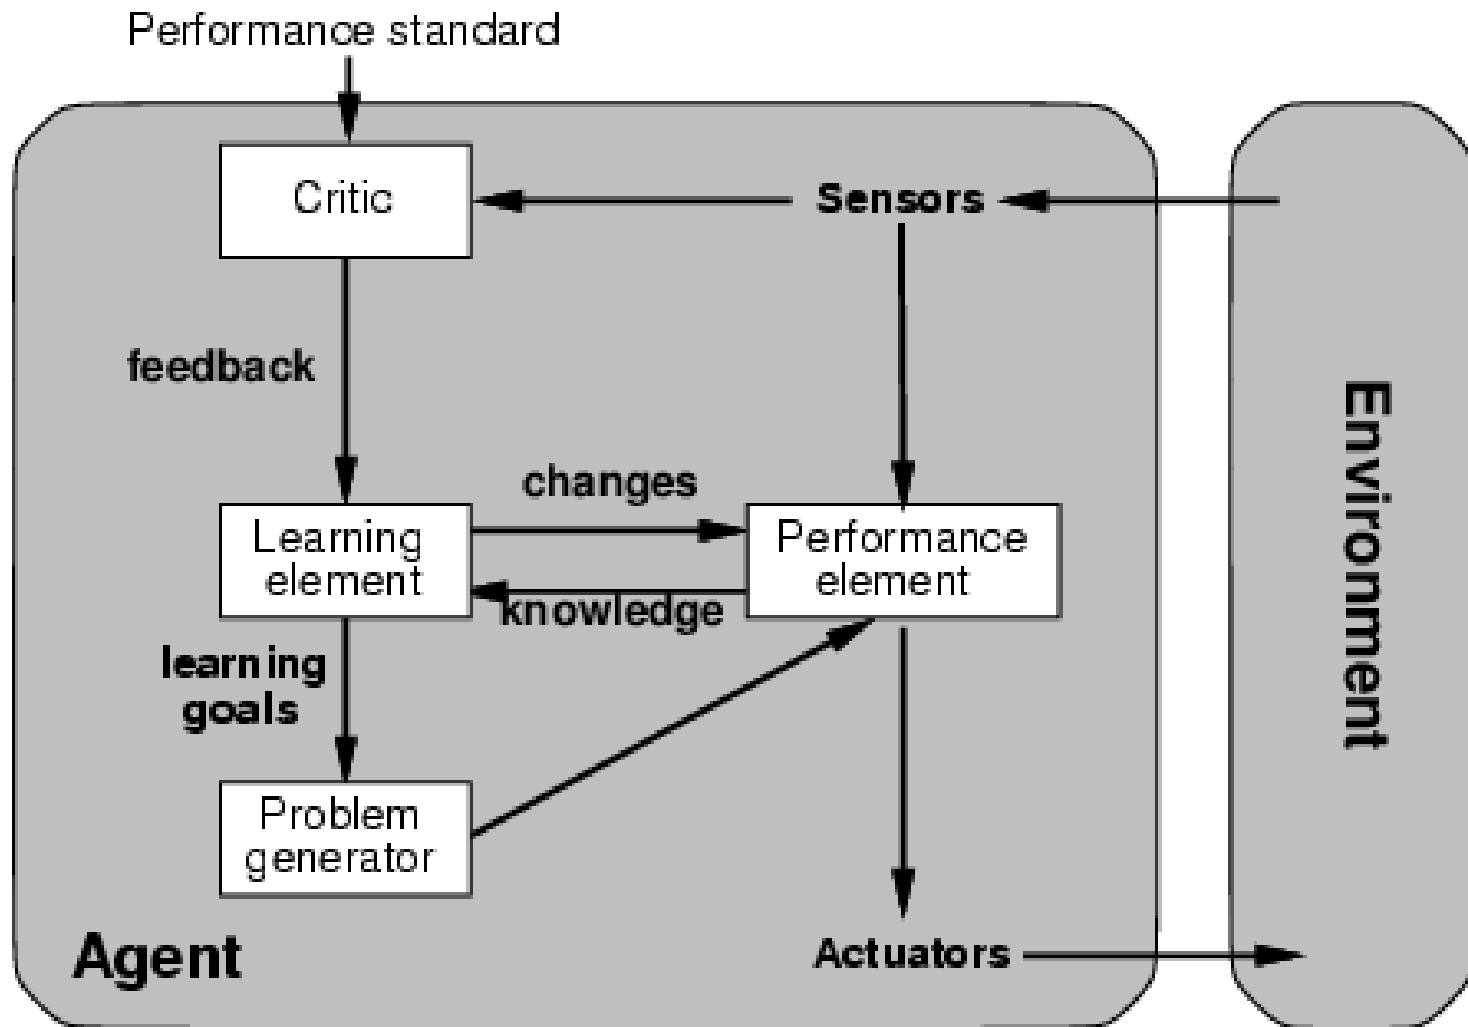


Utility-based agents

The agent uses a utility function to evaluate the desirability of states that could result from each possible action



Learning agents



Questions?

Homework

- 2.4